manvendra@rajpoot.dev
+91 94522 77560

Manvendra Rajpoot

rajpoot.dev
github.com/alzywelzy
linkedin.com/in/alzywelzy

Dear Hiring Manager,

I am writing to apply for the DevOps Engineer role on your team. I come to infrastructure as an engineer who also writes the services running on it, so I care about how a platform is built, shipped, observed and recovered as much as how it is coded. I work daily with Docker and containerised deployments, CI/CD, Linux and Kubernetes concepts, and I would like to make platform work my main focus.

At CloudTechTiq I own the Docker Compose topology for our cloud console: 35 services covering the Django portal, the CloudStack and billing FastAPI microservices, Celery workers and Beat, PostgreSQL, Redis, RabbitMQ, MinIO, NGINX and the observability stack. I wrapped it in one idempotent `make dev` that checks out submodules, rebuilds images only when a Dockerfile or lockfile changed, waits until every application answers, applies pending migrations and prints a health report, so a warm start takes about four seconds. I built GitHub Actions pipelines that run pytest, build multi-stage images and publish them to GHCR for dev, staging and production, with a separate migrations workflow and container health checks behind Gunicorn and NGINX.

I built the observability platform from the ground up: OpenTelemetry instrumentation of FastAPI, SQLAlchemy, Redis, Celery and httpx, routed through an OTel Collector and Grafana Alloy into Prometheus, Loki and Tempo, so a failure anywhere traces back to a single request. I made environments reproducible with database seed snapshots and code-defined reference data, and automated bare-metal provisioning over HPE iLO Redfish with Celery-driven IP allocation. Behind it I designed the services themselves, including the Razorpay and Stripe payment flow and an OAuth2 gateway to Apache CloudStack.

I have also run edge and cloud infrastructure end to end: Cloudflare DNS, Pages and Workers, mail authentication with SPF, DKIM and DMARC at `p=reject`, uptime monitoring, and written decision records and runbooks. I have taken deployments from a GCP VM running a single-node Docker Swarm behind a cloudflared tunnel with automatic image roll-outs, through managed platforms, to serverless, weighing cost, limits and operability at each step. Every migration had a written cutover plan with a check and a rollback per step, and background work runs with leases, heartbeats, retries and dead-letter queues so failures surface instead of disappearing.

I treat infrastructure as something to be explained as well as run: every change has a check, a rollback and a written reason, which is what lets a small team operate it safely. I am keen to deepen my work with Kubernetes and infrastructure as code in a team that runs platforms at a larger scale.

I would welcome the chance to discuss how this experience fits your team, and I am happy to walk through any of the systems above in technical detail. Thank you for your time and consideration.

Sincerely,

Manvendra Rajpoot